

UCT

Polynomial Time

KLAUS SUTNER
CARNEGIE MELLON UNIVERSITY
SPRING 2022



- 1 Tractability
- 2 Weak Reductions
- 3 Sanity Check

Feasible Decision Problems

2

The question arises whether any time complexity class is a good match for our intuitive notion of “feasible computation” or “tractable problem.”

Note that whatever answer we give, we are in a similar situation as with the Church-Turing thesis: since we deal with intuitive notions there cannot be a formal proof – though one can try to amass ample evidence.

However, while there is fairly good agreement about what constitutes a computation, there is much less agreement about what constitutes a feasible computation. If one deals with very large data sets, anything beyond $n \log n$ is not really feasible.

(Deterministic) Polynomial Time

3

Recall our definition of polynomial time solvable problems:

$$\mathbb{P} = \text{TIME}(n^{O(1)}) = \bigcup_{k \geq 0} \text{TIME}(n^k).$$

Let's put up a straw man: let's declare \mathbb{P} to be our stand-in for feasible computation.

But note that \mathbb{P} is only about decision problems, in the world of actual algorithms there is no doubt that function problems are more important.

Types of Problems

4

Here is a fairly natural taxonomy of computational problems.

Decision Problems Return a Yes/No answer.

Counting Problems Determine the size of some object.

Function Problems Calculate a certain function.

Search Problems Select one particular solution.

How about all this other stuff?

Example: Vertex Cover

5

Problem: **Vertex Cover (decision)**

Instance: A ugraph G , a bound k .

Question: Does G have a vertex cover of size k ?

Problem: **Vertex Cover (counting)**

Instance: A ugraph G .

Solution: Find the size of a minimal vertex cover of G .

Problem: **Vertex Cover (function)**

Instance: A ugraph G .

Solution: Find the lex-minimal vertex cover of G .

Problem: **Vertex Cover (search)**

Instance: A ugraph G .

Solution: Find a vertex cover of G of minimal size.

Connections

6

It should be clear that there are lots of connections between these different versions of the problem.

For example, if we can solve the counting/function/search version, then we also can handle the decision version.

The opposite direction is also true, more later.

Feasible Counting/Function/Search Problems 7

To deal with those, we need **transducers** as opposed to **acceptors**. Recall that in a transducers we attach an additional input/output tape to our Turing machines.

Definition

A function is **polynomial time computable** if it can be computed by a polynomial time Turing machine transducer.

The output for counting problems is just a number, but for function and search problems it can be some arbitrary combinatorial object (represented as a string).

Full Disclosure

8

Computing functions, as opposed to just solving decision problems, is beyond any doubt more important in the world of real algorithms. Interestingly, the concept of the class of **polynomial time computable** functions appeared a bit before modern complexity theory:

- von Neumann 1953
- Gödel 1956
- Cobham 1964
- Edmonds 1965

So why don't we bite the bullet and deal with feasible computable functions?

Because it's very hard. Decision problems make enough of a mess, so let's not worry about things that are even more complicated at this point.

Closure

9

We mention just one hugely important property of polynomial time computable functions: they can be composed sequentially without falling out of the class.

Lemma

Polynomial time functions are closed under composition.

Informally, if we have a fast algorithm to take A to B , and another fast algorithm to take B to C , then we also have a fast algorithm to take A to C .

Proof

10

Proof. The short version is: polynomials are closed under substitution.

More precisely, suppose $y = f(x)$ is computable in time at most $p(n)$ where $n = |x|$.

Then $|y| \leq p(n)$ and $z = g(y)$ is computable in time at most $q(p(n))$ for some polynomial q .

Hence we have a polynomial time bound for $z = g(f(x))$. \square

In fact, if $p(n) = O(n^k)$ and $q = O(n^\ell)$, then the composition is $O(n^{k\ell})$ at most.

Polynomials are Great

11

Just to be clear: this result may seem trivial, but it depends crucially on our choice of polynomials as the resource bounds.

Closure fails for, say, exponential time $\text{EXP}_1 = 2^{O(n)}$.

But similar claims do hold for

- linear time (easy),
- logarithmic space (hard, more later).

Two Standard Objections

12

The notation $T_{\mathcal{M}}(n) = O(n^k)$ hides two constants:

$$\exists n_0, c \forall n \geq n_0 (T_{\mathcal{M}}(n) \leq c \cdot n^k)$$

What if these constants are huge? Something like 1000^{1000} ? Note that there are only around 10^{80} particles in the universe[†].

This would render the asymptotic bounds entirely useless for anything resembling feasible computation. It could even wreck physical computation entirely.

[†]Absurdly large n_0 appear with annoying frequency in algorithm analysis, say, fast integer and matrix multiplication.

Standard Answer

13

This objection has merit in principle, but in the RealWorld™ it is pointless: for practical problems it is a matter of experience that the constants are easy to determine and are (almost) always very reasonable.

In fact, we can even compute the constants by writing down the algorithms very carefully in a low-level language like C.

For practical algorithms this is somewhat of a pain, but not really terribly difficult (as long as we are fairly relaxed about the bounds).

Example: Quadratic

14

Consider a code fragment like

```
// P
for i = 1 to n do
  for j = 1 to i do
    blahblahblah
```

Suppose blahblahblah is constant time, a bunch of comparisons and assignments, say.

Clearly, the running time of P is quadratic, $O(n^2)$.

Assembly

15

But if we wanted to, we could write P in assembly,

```
0: 55                push    ebp
1: 89 e5             mov     ebp, esp
3: 83 e4 f0          and     esp, 0xfffffff0
6: 83 ec 10          sub     esp, 0x10
9: c7 04 24 00 00 00 00 mov     DWORD PTR [esp], 0x0
```

Now we can count the number of steps each execution of P takes, likewise for the control structures. The result might be $55n + O(1)$.

But it surely won't be $1000^{1000}n + O(1)$.

Digression: Knuth

16

D. Knuth (everybody bow three times in the direction of Stanford) actually thinks that pure asymptotic bounds are fairly cheesy.

We should never say “this algorithm has complexity $O(n)$ ”, we should explicitly figure out the coefficients of the leading term. Say

$$\dots = 7 \cdot n + O(\log n)$$

This is easy to say for Knuth, who just develops the necessary math on the fly, but quite challenging for ordinary mortals.

The Bibles

17

D. E. Knuth
[The Art of Computer Programming](#)
Addison-Wesley, 1968–now

D. E. Knuth, O. Patashnik, R. Graham
[Concrete Mathematics](#)
Addison-Wesley, 1988


Knuth is very unhappy about people publishing algorithms without ever implementing them (one really should distinguish more carefully between computable function and algorithm).

<http://www.informit.com/articles/article.aspx>

The claim that we can always figure out the multiplicative constant is somewhat of a white lie. There are algorithms in graph theory, based on the Robertson-Seymour theorem for **graph minors** where these constants are utterly unknown.

Graph H is a minor of graph G if H is (isomorphic to) a graph obtained from G by

- edge deletions,
- edge contractions, and
- deletion of isolated vertices.




A collection of finite graphs \mathcal{G} is **closed under minors** if whenever $G \in \mathcal{G}$ and H is a minor of G , then $H \in \mathcal{G}$.

The classical example is the class of planar graphs: every minor of a planar graph is always planar. This produces a famous theorem:

Theorem (Kuratowski-Wagner)

A graph is planar if, and only if, it has no minor K_5 or $K_{3,3}$.



In 1937, K. Wagner conjectured that a similar theorem holds for every minor closed class \mathcal{C} of finite graphs:

Conjecture

Suppose \mathcal{C} is minor closed. There are finitely many obstruction graphs H_1, H_2, \dots, H_r such that G is in \mathcal{C} iff G does not have H_i as a minor for all $i = 1, \dots, r$.

In other words, all anti-chains of graphs (wrt the minor order) are finite.

Note that this yields a decision algorithm: we just check if some given graph G has one of the forbidden minors.

Wagner's conjecture gave rise to the following amazing theorem[†]:

Theorem (Robertson, Seymour)

Every minor-closed family of graphs has a finite obstruction set.

This was proven in a series of 23 papers from 1984 to 2004, an incredible tour de force. The total proof is hundreds of pages long.

The proof has been examined very carefully and is probably correct, but putting it through a theorem prover would not hurt.

[†]This is very closely connected to Higman's theorem about the lack of infinite antichains in the subsequence order.

The proof of the Robertson-Seymour theorem is brutally non-constructive: the finite obstruction set exists in some set-theoretic universe, but we have no way of constructing it.

Worse, often we don't even know its (finite) cardinality.

Does it have any computational content? Well, for planar graphs is certainly does: we could check planarity by looking for minors K_3 and $K_{3,3}$. Needless to say, there are much better planarity testing algorithms.

Theorem To Algorithm

24

Suppose that graph H is **fixed**.

Problem: **Minor Query**
Instance: A ugraph G .
Question: Is H a minor of G ?

There is an algorithm that tests whether H is a minor of G in $O(n^2)$ steps, $n = |G|$.

Note that H has to be fixed, otherwise we could check whether G contains a cycle of length $|G|$ in quadratic time (aka Hamiltonian cycle).

Doom and Disaster

25

But then we can check all H from a finite list of obstruction graphs in quadratic time.

Hence we can check membership in any minor closed class \mathcal{G} in quadratic time.

Alas, there is a glitch: the finite obstruction list is obtained non-constructively; it exists somewhere, and we can prove its existence using sufficiently strong set theory, but we cannot in general determine its elements—and, in fact, not even its cardinality. So we get a quadratic time algorithm, but we cannot bound the multiplicative constant.

Ouch.

Objection Two

26

What if $T_{\mathcal{M}}(n) = O(n^{1000})$?

This is polynomial time, but practically useless. By the time hierarchy theorem, we know that problems exist that can be solved in time $O(n^{1000})$ but essentially not in less time.

But note that these problems are constructed by diagonalization techniques, and are thus entirely artificial; they do not correspond to decent RealWorld™ problems.

Weird Empirical Fact

27

If a natural problem is in \mathbb{P} at all, then it can actually be solved in time $O(n^{10})$ – for some small value of 10.

Take this with ample amounts of salt, but experience has shown so far that there simply are no natural problems where the best known algorithm has running time $O(n^{1000})$.

Alas, no one has any idea why this low-degree principle appears to be true. Note the qualifier “natural”. Everyone understands intuitively what this means, but it seems very difficult to give a formal definition.

Close Call

28

In 2002, Agrawal, Kayal and Saxena published a paper that shows that primality testing is in polynomial time.

Amazingly, the algorithm uses little more than high school arithmetic.

The original algorithm had time complexity $\tilde{O}(n^{12})$, but has since been improved to $\tilde{O}(n^6)$.

Alas, the AKS algorithm seems useless in the RealWorld™, probabilistic algorithms are much superior[†].

1 Tractability

2 Weak Reductions

3 Sanity Check

[†]An indication that \mathbb{P} is probably not quite the right definition of feasible computation.

Scaling Down

30

“Small” classes like \mathbb{P} or EXP don’t play well with our reductions from CRT. For example

$$A \in \text{EXP} \quad \text{implies} \quad A \leq_T \emptyset$$

The problem is that the oracle TM can directly solve A , it has no need to query the oracle. And, \leq_T can eviscerate much more than just EXP .

We need to limit the compute power of the reduction mechanism.

Polynomial Time Turing Reductions

31

Perhaps the most obvious attempt would be to use Turing reductions, but with the additional constraint that the Turing machine must have polynomial running time. This is called a **polynomial time Turing reduction**.

Notation: $A \leq_T^p B$.

Note that this actually makes sense for any kind of problem, not just decision problems. Also, it really captures nicely the idea that problem A is easier than B since we could use a fast algorithm for B to construct a fast algorithm for A .

Free Lunch

32

Just to be clear: in a computation with oracle B we only charge for the steps taken by the ordinary part of the Turing machine. For example, it costs to write a query for the oracle on the tape.

But the answer provided by the oracle appears in just one step, we do not care about the internals of the oracle.

Checking

33

Proposition

Polynomial time Turing reducibility is a preorder (reflexive and transitive).

For transitivity, this works since polynomials are closed under substitution. Hence we can form equivalence classes as usual, the **polynomial time Turing degrees**. We won't pursue this idea here.

Compatibility

34

Proposition


$B \in \mathbb{P}$ and $A \leq_T^p B$ implies $A \in \mathbb{P}$.

Here we can simply replace the oracle by a real polynomial algorithm.

Since polynomials are closed under substitution, this will produce a polynomial time algorithm.

Recall: Vertex Cover

35



Finding minimal vertex covers is quite messy, even for moderate size graphs. In fact, even verifying a solution is hard.

Comparisons

36

Recall the 4 versions: decision, counting, function and search.

Intuitively, the function version is the hardest: once we have a lex-minimal cover, all other problems are trivial. Clearly, they are all polynomial Turing reducible to the function version.

The counting and decision versions are also Turing reducible to the search version.

Lastly, the decision version is Turing reducible to counting.

$$\text{Decision} \leq_T^p \text{Counting} \leq_T^p \text{Search} \leq_T^p \text{Function}$$

The Other Way

37

Proposition

The function version of Vertex Cover is polynomial time Turing reducible to the decision version.

This requires a little argument. Let n be the number of vertices in G , say, $V = \{v_1, \dots, v_n\}$ in lexicographic order.

- First, do a binary search to find the size k_0 of a minimal cover, using the oracle for the decision problem.
- Then find the least vertex v such that $G - v$ has a cover of size $k - 1$. Place v into C and remove it from G .
- Recurse.

This is polynomial time, even using Turing machines.

Justification

38

The last proposition is important: it justifies our narrow focus on decision problems: up to a polynomial factor, they are no worse than the fancier, more applicable versions. They all live in the same polynomial time Turing degree.

Of course, this does not help much when we are trying to find super-efficient algorithms. Typical example: Tarjan's strongly connected component algorithm, a form of DFS on steroids.

But, if all we want is to make sure that things can be done in polynomial time, the decision version is often good enough. Ditto for showing that things cannot happen in polynomial time.

Still No Good

39

Unfortunately, while polynomial Turing reductions are compatible with \mathbb{P} , they are still too brutish. To wit:

$$A \in \mathbb{P} \quad \text{implies} \quad A \leq_T^p \emptyset$$

The problem again is that the oracle TM can directly solve A , it has no need to query the oracle.

How about the other reductions (many-one in particular) that we considered in the CRT part?

Weak Reductions

40

Many-one reducibility is based on a computable function translating from one problem to another. We'll just have to constrain the function a bit.

Definition

$A \subseteq \Sigma^*$ is **polynomial time (many-one) reducible** to $B \subseteq \Sigma^*$ if there is a polynomial time computable function f such that

$$x \in A \iff f(x) \in B.$$

Notation: $A \leq_m^p B$.

Properties

41

Proposition

Polynomial time many-one reducibility is a preorder.

Proof.

Reflexivity is trivial. For transitivity, consider a polynomial time reduction f from A to B and g from B to C .

Obviously $h(x) = g(f(x))$ is a reduction from A to C .

h is still polynomial time computable since polynomials are closed under substitution.

□

Better Mousetrap?

42

Proposition

\leq_m^p is compatible wrto \mathbb{P} : if B is in \mathbb{P} and $A \leq_m^p B$, then A is also in \mathbb{P} .

This is clear: we can replace the oracle B by a polynomial time algorithm.

But is \leq_m^p really a useful reduction for \mathbb{P} ?

We could use the mediating function to “reduce” a problem in $\text{TIME}(n^{1000})$ to $\text{TIME}(n)$.

Again: we really want the reductions to be lightweight, the oracle should be the place where the heavy lifting takes place.

Log-Space Reductions

43

As it turns out, a better reduction results from restricting f even more: we'll insist that f can be computed in **logarithmic space**. More details about space complexity later, for the time being use your intuition.

Definition


$A \subseteq \Sigma^*$ is **log-space reducible** to $B \subseteq \Sigma^*$ if there is a log-space computable function f such that

$$x \in A \iff f(x) \in B.$$

Notation: $A \leq_\ell B$.

Recall: Transducer

44



Only the work tape counts. We want it to be logarithmic in the size of the input.

But Why?

45

Clearly $A \leq_\ell B$ implies $A \leq_m^p B$.

But the mediating function f is now much more constrained. Think about graph problems where $[n]$ is the vertex set.

We cannot remember an arbitrary subset $S \subseteq [n]$.

Unfortunately, this wrecks many graph algorithms such as DFS that require linear memory and nicely forces the oracle to do most of the work.

Example: Neighborhood Traversals

46

Here is a typical code fragment in a graph algorithm:

```
foreach v in V do
  foreach (v,u) in E do
    blahblahblah
```

This idiom appears in lots of graph algorithms: look at the neighborhoods of all vertices.

Good news: it can be handled in logarithmic space: all we need is a few vertices in $[n]$ (these are written in binary), we can traverse all the edges (v, u) in any standard representation of a graph (this works fine for both adjacency matrices or adjacency lists, though they may produce different running times).

Structuring Polynomial Time

47

Unlike polynomial-time reductions, log-space reductions can be used to study the fine-structure of \mathbb{P} .

Definition (\mathbb{P} -Completeness)

A language B is **\mathbb{P} -hard** if for every language A in \mathbb{P} : $A \leq_\ell B$. It is **\mathbb{P} -complete** if it is \mathbb{P} -hard and in \mathbb{P} .

Building a \mathbb{P} -hard set is easy: take an enumeration (\mathcal{M}_e) of all polynomial time TMs, and build the analogue to the Halting set:

$$K_p = \{ e \# x \mid \mathcal{M}_e(x) \text{ accepts} \}$$

Of course, there is no reason why K_p should be in \mathbb{P} .

To obtain (\mathcal{M}_e) , we start with an arbitrary effective enumeration of all Turing machines and filter out the good machines. Let's ignore the problem that we cannot check whether the machines are total.

We would still need to run a poly-time tester that checks that \mathcal{M}_e runs in polynomial time.

Unfortunately, this tester does not exist either: it is undecidable whether a machine runs in polynomial time. Intuitively we need to check

$$\exists k \forall n, x (|x| = n \Rightarrow T_{\mathcal{M}}(x) \leq n^k + k)$$

This looks completely hopeless, something like Σ_2 .

But we can construct a new machine \mathcal{M}_e as follows. As always, assume that an enumeration (N_e) of all machines has infinitely many repetitions.

- attach a clock to N_e , and
- stop the computation after at most $n^e + e$ steps;
- return the same output as $N_e(x)$ if the computation has converged,
- otherwise return 0 (or some other default value).

Note that \mathcal{M}_e is total by definition.

Furthermore, if some total machine \mathcal{M} runs in polynomial time, then there is an index e such that $\mathcal{M}(x) \simeq \mathcal{M}_e(x)$ for all x .

If \mathcal{M} is partial, we get the same behavior on the support.

Problem: **Circuit Value Problem (CVP)**
Instance: A Boolean circuit C , input value x .
Question: Check if C evaluates to true on input x .

Obviously CVP is solvable in polynomial time (even linear time given a halfway reasonable representation).

There are several versions of CVP, here is a particularly simple one: compute the value of the “last” variable X_m given

$$\begin{aligned} X_0 &= 0, & X_1 &= 1 \\ X_i &= X_{L_i} \diamond_i X_{R_i}, & i &= 2, \dots, m \end{aligned}$$

where $\diamond_i = \wedge, \vee$ and $L_i, R_i < i$.

Theorem (Ladner 1975)
The Circuit Value Problem is \mathbb{P} -complete.

Sketch of proof. [†] For hardness consider any language A accepted by a polynomial time Turing machine \mathcal{M} .

We can encode the computation of the Turing machine \mathcal{M} on x as a polynomial size circuit (polynomial in $n = |x|$): use lots of Boolean variables to express tape contents, head position and state.

Constructing this circuit only requires “local” memory; for example we need $O(\log n)$ bits to store a position on the tape.

The circuit evaluates to true iff the machine accepts.

□

[†]We will provide much more detail in the proof of the Cook-Levin theorem.

A long list of \mathbb{P} -complete problems is known, though they tend to be a bit strange.

For example, consider an undirected graph $G = \langle V, E \rangle$ where E is partitioned as $E = E_0 \cup E_1$.

Question: Given a start vertex s and a target t , does t get visited along an edge in E_0 if the breadth-first search can only use edges in E_0 at even and E_1 at odd levels?

Theorem
ABFS is \mathbb{P} -complete (wrt to logspace reductions).

- 1 Tractability
- 2 Weak Reductions
- 3 Sanity Check

In any standard algorithm text you will find a statement along the lines of

Proposition

DFS runs in linear time.

What does that mean in our framework?

First, the size of the input, a ugraph $G = \langle V, E \rangle$ is usually given as $n + e$ where $n = |V|$ and $e = |E|$. This is justified by the standard adjacency list representation, and uniform cost function: in the RealWorld™, no one deals with graphs larger than $2^{64} \approx 1.85 \times 10^{19}$, much less 2^{128} .

If we were to implement DFS on a Turing machine, there would obviously be some slow-down. But still, there is a low-degree polynomial p so that the Turing machine runs in time $p(n + e)$.

It is great fun to figure out how small p can be made, but ultimately irrelevant: we want **lower bounds**, in particular statements like

Problem such-and-such **cannot** be solved in polynomial time.

A polynomial increase doesn't matter in this case: we might as well think about the computation in a luxurious RAM or a C program. If the Turing machine does not run in polynomial time, these won't either.

This makes it much easier to reason about the problem.

Given a digraph and two vertices s and t , we can ask whether there is a path from s to t .

The brute-force approach would enumerate all simple paths starting at s , and check if any one of them hits t . Alas, this is exponential.

But we can simply run DFS (or BFS) starting at s , and get the answer in linear time.

Given a digraph whose edges are labeled by a positive cost, one can ask: what is the shortest path from s to t ?

Again, brute-force would require enumeration of all simple paths and require exponential time.

But there are better (and more complicated) algorithms such as Dijkstra or Bellman-Ford that work in polynomial time.

Just this once, let us make sure that a Turing machine could handle, say, arrays in polynomial time, and even at fairly low degree.

We want a reasonable Turing implementation of **arrays** of integers: a mutating data structure to store a (fixed-length) sequence of integers. The key operations are read and write:

```
x := A[i];      \\ read
A[i] := x;      \\ write
```

To simplify matters, we will use three tapes on the TM to represent a single array: one for the actual array, one for the index i , and one for the value x .

The array tape looks like

$$\#a_0\#a_1\#\dots\#a_{n-1}\#$$

where the numbers a_i are written in binary, and $\#$ is a special separator symbol.

For a read operation, we use the index i to find the i th separator, searching from the left. Then we copy the following binary number to the value tape.

For a write operation, we again search for the i th block. We then copy x over from the other tape, shifting the remainder of the tape contents according to the number of bits in x (compared to the previous number of bits).

Alternatively, we could determine the number k of bits of the largest number ever stored, and reserve k bits in each block right from the start.

Of course, this wrecks the crucial property of an array, constant time read/write access. But: It only takes $O(nk)$ steps to simulate the ops.

More generally, when we implement some graph algorithm such as Dijkstra's shortest path on a Turing machine, the whole computation slows down, but only by a polynomial amount. In particular, if our pseudo-code runs in polynomial time, then the TM also runs in polynomial time. So we actually have a proof (sketch) that the shortest path problem is in \mathbb{P} .

Again: we are losing all fine-grained information, but this is fine if one is mostly interested in separating feasible from infeasible. The key accomplishment of Dijkstra's algorithm is to avoid an exponential search, the TM still reflects this.

Given a ugraph, we can ask: is there a simple path that has length at least k ?

Certainly brute-force enumeration of all simple paths would dismantle this problem, and would require exponential time.

But this time no one knows a computational shortcut, in particular no polynomial time algorithm for this problem is known.

The problem gets nasty when $k = n - 1$: we are asking for a simple path using all vertices. This seems to be fundamentally different from anything like a shortest path.

As we will see, there are lots of combinatorial problems like the last one that are easy to solve with an exponential search, everything else is polynomial.

Naturally, one would like to establish a lower bound, say, every algorithm requires 2^n steps (at least for infinitely many inputs).

Alas, these lower bounds seem be extremely elusive, to the point of constituting a major open problem in math, not just the theory of algorithms.

We already talked about one other such problem: Vertex Cover.

We are given a ugraph G with n vertices and a bound $k \leq n$, and want to know whether G admits a VC of size k .

Obviously we can perform a brute force check on all subsets of size k , but there are $\binom{n}{k}$ such sets—and that number is **not** polynomial in n for variable k (though everything is fine for fixed k).

So, on the face of it, it is far from clear that VC can be handled in polynomial time.